

Automated 3D Stenosis Quantification and Visualization Workflow for Coronary Artery Disease Clinical Decision Support

Adrià Cortés Cugat



Universitat
Pompeu Fabra
Barcelona

A la meva família, que ha estat al meu costat en tot moment, donant-me l'empenta i el suport incondicional que necessitava, sobretot quan més em costava creure en mi mateix. I als meus amics, per ser la força que m'ha acompanyat dia a dia. Aquest projecte no marca només el final del grau, sinó el tancament d'una etapa vital molt important que m'ha canviat com a persona i l'inici d'un nou capítol.

A l'Hospital de Sant Pau i al Dimension Lab, per obrir-me les portes i ensenyar-me que el més bonic d'aquesta professió no són els algorismes ni els números, sinó veure com allò que aprens pot servir per ajudar les persones i fer el bé. Sentir que aquest esforç pot tenir un impacte real i útil en la vida dels altres és la recompensa que dona sentit a tot el camí.

Acknowledgements

First and foremost, I would like to express my sincere gratitude to Oscar Camara Rey, my thesis director at Universitat Pompeu Fabra, for the trust he placed in me from the very first day, his remarkable ability to align my technical interests with real clinical needs, and for opening the doors to a medical sector that was completely new to me.

I also want to thank César Acebes Pinilla, my supervisor at Hospital de la Santa Creu i Sant Pau, for welcoming me into the Dimension Lab and guiding me through a real coronary artery disease diagnostic environment. Furthermore, my sincere thanks go to the expert cardiologists and medical doctors who took the time to review my work, provided invaluable feedback, and explained their daily workflows, allowing me to adapt this project to their actual needs.

I further acknowledge the prior work of the students and researchers within this clinical ecosystem, upon which much of the technical framework of this thesis is built.

Finally, I want to thank my family and my friends, both from the Mathematical Engineering in Data Science degree and other engineering programs, who have accompanied and supported me throughout my entire university journey at UPF.

Abstract

Coronary artery disease (CAD) assessment from coronary computed tomography angiography (CCTA) remains largely manual and time-consuming in clinical practice. At Hospital de la Santa Creu i Sant Pau, expert review of a single study can require between 30 minutes and two hours depending on case complexity. This bachelor's thesis presents an automated, end-to-end computational pipeline that transforms segmented coronary anatomy into quantitative stenosis metrics, CAD-RADS 2.0-oriented reporting support, and an interactive clinical visualization dashboard. The system is organized into four modular blocks: automated anatomy extraction, geometric stenosis quantification, segment labeling with CAD-RADS scoring, and a Streamlit-based decision-support interface. Designed as a transparent in-house framework rather than a proprietary black box, the pipeline aims to reduce reporting burden, improve reproducibility, and assist patient prioritization while preserving final clinical authority. Validation combines ASOCA and MACS-18 cohorts with synthetic phantoms for controlled geometric verification.

Resum

L'avaluació de la malaltia coronària (MAC) a partir de la tomografia computada coronària (TCCC) continua sent, en la pràctica clínica, un procés manual i molt costós en temps. A l'Hospital de la Santa Creu i Sant Pau, la revisió experta d'un estudi pot requerir entre 30 minuts i dues hores segons la complexitat del cas. Aquest treball de fi de grau presenta un pipeline computacional automatitzat d'extrem a extrem que transforma l'anatomia coronària segmentada en mètriques quantitatives d'estenosi, suport orientat al CAD-RADS 2.0 i un tauler clínic interactiu de visualització. El sistema s'organitza en quatre blocs modulars: extracció automatitzada de l'anatomia, quantificació geomètrica de l'estenosi, etiquetatge segmentari amb puntuació CAD-RADS i una interfície Streamlit de suport a la decisió. Dissenyat com a marc transparent i intern, el pipeline pretén reduir la càrrega de informes, millorar la reproduïbilitat i assistir la priorització de pacients sense substituir l'autoritat clínica final. La validació combina les cohorts ASOCA i MACS-18 amb fantasmes sintètics per a la verificació geomètrica controlada.

Resumen

La evaluación de la enfermedad coronaria (EC) a partir de la angiografía por tomografía computarizada coronaria (ATCC) sigue siendo, en la práctica clínica, un proceso manual y muy costoso en tiempo. En el Hospital de la Santa Creu i Sant Pau, la revisión experta de un estudio puede requerir entre 30 minutos y dos horas según la complejidad del caso. Este trabajo de fin de grado presenta un *pipeline* computacional automatizado de extremo a extremo que transforma la anatomía coronaria segmentada en métricas cuantitativas de estenosis, soporte orientado al CAD-RADS 2.0 y un panel clínico interactivo de visualización. El sistema se organiza en cuatro bloques modulares: extracción automatizada de la anatomía, cuantificación geométrica de la estenosis, etiquetado segmentario con puntuación CAD-RADS e interfaz Streamlit de apoyo a la decisión. Diseñado como un marco transparente e interno, el *pipeline* pretende reducir la carga de informes, mejorar la reproducibilidad y asistir la priorización de pacientes sin sustituir la autoridad clínica final. La validación combina las cohortes ASOCA y MACS-18 con fantasmas sintéticos para la verificación geométrica controlada.

Preface

This document presents the bachelor’s thesis completed during the 2025–2026 academic year within the Bachelor’s Degree in Mathematical Engineering in Data Science at Universitat Pompeu Fabra, in collaboration with the Dimension Lab at Hospital de la Santa Creu i Sant Pau.

The main body of the thesis (Introduction through Conclusions) is written in English for consistency with the scientific literature in the field.

Use of Generative AI Disclosure

The Author, Adrià Cortés Cugat, discloses that the writing of this document “Automated 3D Stenosis Quantification and Visualization Workflow for Coronary Artery Disease Clinical Decision Support” involved the use of Cursor AI and Google Gemini AI, for the purpose of: language review; text editing; technical writing support; $\LaTeX$ integration and formatting; software development assistance; code debugging and optimization; and pipeline documentation support. The Author warrants, however, that the contents of this thesis are a genuine human creation, as the results obtained from Cursor AI and Google Gemini AI have been duly reviewed and edited by Adrià Cortés Cugat. In accordance with university guidelines on the use of AI in assessable work, the technical and clinical substance of the main body, especially the Introduction, Methodology, and Conclusions, has been written, validated, and revised by the Author with the support of the thesis director and supervisor.

Information source disclosure. The Author declares that all cited sources have been appropriately reviewed and validated by Adrià Cortés Cugat to support the claims made in this work, ensuring they are sufficiently contrasted and reliable.

Note on Figures and Data

The diagram figures that illustrate clinical and reporting workflows have been entirely designed and created by the Author. Where other figures build upon internal hospital materials, they are explicitly cited in the corresponding figure captions. Clinical cohort data and pipeline results are documented in the corresponding chapters and, where appropriate, in the supporting appendix.

Contents

List of Figures	IX
List of Tables	X
1. INTRODUCTION	1
1.1. Clinical Context and Motivation	1
1.1.1. Anatomy of Coronary Arteries	1
1.1.2. Coronary Artery Disease	2
1.1.3. Stenosis Quantification: A Key Indicator of CAD Severity	3
1.1.4. CAD-RADS 2.0 and the Need for Decision Support	3
1.2. Clinical Automation Challenges at Hospital de la Santa Creu i Sant Pau	4
1.2.1. Current CAD Diagnostic Workflow and Associated Limitations	4
1.2.2. The Bottleneck: Manual Image Analysis and Reporting	4
1.3. State of the Art: Automation of CAD Diagnosis and Visualization	5
1.4. Project Scope and Contributions	6
1.5. Objectives	7
2. METHODOLOGY	8
2.1. System Overview and Core Automated Workflow	8
2.1.1. End-to-End Data Flow	8
2.2. ASOCA and MACS-18: Coronary Segmentation and Labeling Data	9
2.2.1. Segmentation and Label Volumes	9
2.2.2. Datasets: ASOCA and MACS-18	10
2.3. Block 1: 3D Artery Anatomy Extraction	10
2.4. Block 2: Geometric Stenosis Quantification	12
2.5. Block 3: Artery Labeling and CAD-RADS 2.0 Scoring	15
2.5.1. Anatomical Context: AHA 17-Segment Model	15
2.5.2. Segment-Level Stenosis Aggregation	16
2.5.3. Patient-Level CAD-RADS 2.0 Classification	17
2.6. Block 4: Patient Visualization Dashboard	18
2.7. Experimental Design and Validation Framework	19
2.7.1. Real-Data Evaluation Metrics and Acceptance Criteria	20
2.7.2. Workflow Validation on Real Datasets: ASOCA	21
2.7.3. Workflow Validation on Real Datasets: MACS-18	21
2.7.4. Geometry and Quantification Verification: Synthetic Data	21
2.7.5. User Experience and Interface Usability Evaluation	22
3. RESULTS	23
3.1. Pipeline Execution: ASOCA vs. MACS-18	23
3.1.1. Execution and Endpoint Topology	23
3.1.2. Ostium and Centerline Quality	24
3.1.3. Qualitative Geometry	25

3.1.4. Runtime	26
3.2. Geometric and Algorithmic Accuracy via Synthetic Data	27
3.2.1. Phantom Overview	27
3.2.2. Synthetic_1: Healthy Control	27
3.2.3. Synthetic_2: Stenosis Control	28
3.3. Coronary Artery Dashboard Visualization Output	30
3.4. User Experience and Interface Usability Evaluation	32
4. DISCUSSION	34
4.1. Analysis of the Results	34
4.2. Technical and Anatomical Limitations	34
4.3. Future Work	34
5. CONCLUSIONS	35
A. SUPPORTING MATERIAL	39
A.1. System Usability Scale Questionnaire	39

List of Figures

1.1.	Anterior view of the human coronary tree. Coronary arteries are highlighted in red, and other cardiac structures in blue. Adapted from Lynch et al. [25].	1
1.2.	Coronary arteries narrowed by atherosclerotic plaque. Adapted from Kaiser Permanente [20].	2
1.3.	Current CAD diagnostic workflow at Hospital de Sant Pau. The primary operational bottleneck (image assessment and reporting) is highlighted in red. Adapted from Ferrer Beltran [15].	4
1.4.	Detailed representation of the image analysis and reporting workflow at Hospital de Sant Pau. Adapted from Acebes Pinilla [1].	5
1.5.	Contextual overview of the proposed automated diagnostic framework at Hospital de la Santa Creu i Sant Pau. The contributions of this thesis (green dashed box) bridge the gap between existing hospital data systems (blue) and the final clinical report. Adapted from Acebes Pinilla [1].	7
2.1.	End-to-end automated workflow for a single patient case. Each block constitutes a self-contained processing stage; arrows indicate mandatory sequential dependencies and the principal data passed between stages.	8
2.2.	Block 1 hybrid workflow: spatial separation of the right and left coronary trees, automated seed placement on each tree, and extraction of centerline geometry with branch packaging.	10
2.3.	Block 2 workflow: orthogonal cross-sectional area extraction on each artery centerline and surface mesh, tabular integration, proximal/distal reference estimation, and %AS quantification.	12
2.4.	Block 3 workflow: segment aggregation based on the AHA 17-segment model and automated CAD-RADS 2.0 scoring.	15
2.5.	Block 4 user interface layout and interaction flow: from global 3D anatomy and clinical summary to segment-level and branch-level metric validation.	18
3.1.	Distribution of detected endpoints per patient for ASOCA and MACS-18 cohorts.	24
3.2.	Grouped comparison of ostium accuracy and centerline precision between ASOCA and MACS-18.	25
3.3.	Representative centerline extraction comparison between ASOCA and MACS-18 masks for the same patient case.	25
3.4.	Representative smoothed surface meshes derived from ASOCA and MACS-18 segmentations for the same patient case.	26
3.5.	Per-block and total pipeline runtimes for ASOCA and MACS-18 cohorts.	26
3.6.	Synthetic healthy and stenosis tube phantoms with extracted centerlines (Equation (2.4)).	27
3.7.	Longitudinal area and %AS profiles along the healthy synthetic tube.	27
3.8.	Three-dimensional area and %AS colormaps for the healthy synthetic tube.	28
3.9.	Longitudinal area and %AS profiles along the stenosis synthetic tube.	29
3.10.	Three-dimensional area and %AS colormaps for the stenosis synthetic tube.	29

3.11. Dashboard Module 1: interactive global RCA and LCA three-dimensional renderings with metric colormaps.	30
3.12. Dashboard Module 2: patient clinical summary with automated CAD-RADS 2.0 category, SIS, and territory rationale.	30
3.13. Dashboard Module 3: segment-level three-dimensional map and synchronized two-dimensional area and %AS bar plots.	31
3.14. Dashboard Module 3 (continued): segment selection and annotated peak stenosis with reference markers.	31
3.15. Dashboard Module 4: branch-level three-dimensional path highlighting and longitudinal metric profiles.	32
3.16. Dashboard Module 4 (continued): full-branch area and %AS profiles in proximal-to-distal order.	32

List of Tables

2.1. Overview of the four pipeline blocks: processing role, primary inputs, and principal outputs per patient case.	9
2.2. AHA 17-segment dictionary used for anatomical labeling and CAD-RADS territory grouping. Adapted from Austen et al. [5].	16
2.3. CAD-RADS 2.0 luminal stenosis categories applied to maximum segment-level %AS. Adapted from Cury et al. [14].	17
2.4. Major coronary territories used for CAD-RADS 2.0 triage rules in this pipeline (segment IDs follow the AHA 17-segment model). Adapted from Austen et al. and Cury et al. [5, 14].	17
2.5. Real-data evaluation metrics used in Sections 2.7.2 and 2.7.3.	20
2.6. Acceptance bands for real-data pipeline metrics.	20
2.7. Acceptance bands for synthetic phantom validation.	22
2.8. SUS score interpretation bands (Bangor adjective scale).	22
3.1. Endpoint summary statistics per cohort and subgroup. MPE denotes mean endpoints per patient; EISR is reported for ASOCA only (Table 2.5).	23
3.2. Ostium accuracy (OISR) and centerline precision (CESR) by cohort and subgroup (%). .	24
3.3. Synthetic_1 validation metrics ($W = 10$ mm).	27
3.4. Synthetic_2 validation metrics ($W = 10$ mm).	28
3.5. Raw Likert responses (1 = strongly disagree, 5 = strongly agree) per SUS item.	33
3.6. SUS score summary per participant and cohort statistics.	33

Chapter 1

INTRODUCTION

1.1. Clinical Context and Motivation

1.1.1. Anatomy of Coronary Arteries

The human heart is the central organ of the cardiovascular system and requires a continuous supply of oxygen and nutrients to function properly. This supply is provided by the coronary artery tree, a specialized vascular network that delivers blood directly to the heart muscle, known as the myocardium. The coronary arteries originate from the root of the aorta and surround the surface of the heart in a crown-like arrangement, from which the term *coronary artery* is derived [13]. As illustrated in Figure 1.1, the coronary arterial system is mainly composed of two vessels: the right coronary artery (RCA) and the left coronary artery (LCA).




Figure 1.1: Anterior view of the human coronary tree. Coronary arteries are highlighted in red, and other cardiac structures in blue. Adapted from Lynch et al. [25].

The RCA primarily supplies blood to the right atrium and right ventricle. In contrast, the LCA quickly divides into two major branches: the left anterior descending (LAD) artery and the circumflex (LCx), which together supply the majority of the left ventricle and the interventricular septum [13]. From these primary vessels, a network of smaller branches extends deep into the myocardium to ensure regional blood delivery across the entire heart.

A critical characteristic of the coronary circulation is its limited collateral connectivity. Unlike other vascular systems in the human body, coronary arteries generally lack sufficient alternative pathways to

redirect blood flow when an obstruction occurs. Consequently, any narrowing or blockage in a proximal arterial segment directly compromises blood flow to all downstream regions supplied by that vessel [28].

Additionally, coronary anatomy presents significant inter-patient variability in terms of vessel size, branching structure, and coronary dominance. Dominance is defined by which artery gives rise to the posterior descending artery (PDA) and can be classified as right-dominant, left-dominant, or codominant. Right-dominant circulation, in which the PDA originates from the RCA, is the most common configuration, followed by left dominance and codominance [32]. Understanding this complex anatomical layout is fundamental, as it provides the basis for locating lesions, quantifying stenosis severity, and stratifying risk in coronary artery disease (CAD) diagnosis.

1.1.2. Coronary Artery Disease

Coronary artery disease (CAD) remains one of the leading causes of morbidity and mortality worldwide [31]. It is characterized by the progressive accumulation of atherosclerotic plaque within the walls of the coronary arteries, leading to a narrowing of the vessel lumen known as stenosis. This narrowing restricts blood flow to the myocardium and, if left untreated, can result in myocardial ischemia or infarction [23, 26]. Figure 1.2 illustrates this pathological mechanism, contrasting a normal artery with one obstructed by plaque buildup.




Figure 1.2: Coronary arteries narrowed by atherosclerotic plaque. Adapted from Kaiser Permanente [20].

Early and accurate diagnosis is essential for effective risk stratification and patient management. In current clinical practice, coronary computed tomography angiography (CCTA) is the standard non-invasive imaging technique for CAD evaluation [23]. CCTA provides detailed 3D reconstructions of the coronary tree, allowing clinicians to detect plaques, analyze their spatial location, and estimate the degree of luminal narrowing.

However, despite the high resolution of CCTA, the diagnostic workflow remains highly time-consuming and heavily dependent on manual and visual assessment. For instance, at Hospital de la Santa Creu i Sant Pau, reporting non-pathological or low-complexity cases requires approximately 30 minutes, while complex evaluations can take up to 2 hours [1]. This significant clinical burden drives the need for automated, data-driven tools capable of quantifying stenosis efficiently and supporting clinical decision-making, while fully preserving expert interpretability and clinical control.

1.1.3. Stenosis Quantification: A Key Indicator of CAD Severity

Among the various indicators used to assess coronary artery disease, stenosis severity plays a central role in clinical diagnosis and risk stratification. Stenosis represents the reduction of the vessel lumen caused by atherosclerotic plaque accumulation, a condition directly associated with an increased risk of myocardial ischemia and adverse cardiac events [16, 23]. In clinical practice, this severity is quantified using geometric measurements such as the minimal lumen area (MLA) or percentage diameter stenosis (%DS). These metrics evaluate the degree of narrowing by comparing the most constricted point of a lesion with a reference baseline, which is typically estimated from adjacent, relatively healthy vessel segments [16, 18].

Although these definitions are conceptually simple, accurate stenosis quantification presents significant clinical challenges. Coronary arteries exhibit complex, tortuous geometries, and diffuse atherosclerosis often makes it difficult to identify truly healthy reference regions. In addition, limitations in image resolution, segmentation errors, and noise in CCTA-derived data can highly influence measurements, leading to variability in stenosis estimates [18, 21].

In routine clinical practice, stenosis assessment is often performed visually or using semi-automatic tools provided by proprietary software [1, 14]. Although these methods are widely used, they depend heavily on the clinician’s interpretation and can vary between observers, which reduces consistency and reproducibility [18].

Consequently, there is a strong clinical need to develop robust, transparent, and reproducible quantification methods. Automating these geometric measurements has the potential to standardize stenosis evaluation, reduce reporting time, and provide reliable quantitative data essential for downstream clinical tasks, such as CAD-RADS 2.0 assignment and patient prioritization.

1.1.4. CAD-RADS 2.0 and the Need for Decision Support

To standardize the reporting of coronary disease and facilitate clinical communication, the Coronary Artery Disease-Reporting and Data System (CAD-RADS 2.0) was established [14]. This structured framework categorizes patients into distinct risk scores ranging from 0 to 5, primarily based on the most severe stenosis measured via CCTA. Each categorical score reflects the overall disease burden and provides specific clinical recommendations, offering a standardized pathway for subsequent patient management [14].

While CAD-RADS 2.0 significantly improves reporting consistency across institutions, its assignment still relies heavily on the manual interpretation of stenosis severity across multiple vessels and segments. This process requires clinicians to continuously synthesize quantitative geometric measurements with complex anatomical context. Consequently, evaluating these scans becomes highly cognitively demanding, particularly in time-constrained clinical environments [1]. In practice, manual scoring remains susceptible to inter-observer variability, especially in borderline cases or in patients presenting with multiple moderate lesions. Such discrepancies can lead to inconsistent diagnostic classifications, ultimately impacting downstream clinical decisions, further testing, and patient prioritization [1, 18].

To address these limitations, there is a growing clinical imperative to develop automated decision support tools. These solutions can span from rule-based algorithms mapping quantitative stenosis measurements to machine learning approaches that leverage structured imaging features. Crucially, these computational tools aim to provide objective, reproducible, and transparent suggestions to assist radiologists and cardiologists. By streamlining the CAD-RADS assignment process, these systems reduce reporting times and variability while strictly preserving expert clinical judgment and final decision authority.

1.2. Clinical Automation Challenges at Hospital de la Santa Creu i Sant Pau

1.2.1. Current CAD Diagnostic Workflow and Associated Limitations

The current diagnostic workflow for Coronary Artery Disease (CAD) at Hospital de la Santa Creu i Sant Pau involves a complex and coordinated series of steps, managed primarily by the specialized cardiac imaging unit (Unitat d'Imatge i Funció Cardíaca, UIFC). As illustrated in Figure 1.3, the clinical pathway begins when a patient presenting with symptoms such as chest pain is scheduled for a Coronary Computed Tomography Angiography (CCTA). Once the scan is acquired, the images are stored in the hospital's *Picture Archiving and Communication System* (PACS) for further evaluation.




Figure 1.3: Current CAD diagnostic workflow at Hospital de Sant Pau. The primary operational bottleneck (image assessment and reporting) is highlighted in red. Adapted from Ferrer Beltran [15].

The core of this diagnostic pipeline relies heavily on the visual and manual analysis of these scans by expert cardiologists and radiologists to extract clinical metrics. These findings are then compiled into a clinical report, which the referring clinician evaluates to determine the appropriate subsequent medical actions, ranging from discharging a healthy patient to programming further invasive tests.

While the overall infrastructure is well-established, the central phase of image analysis and reporting (highlighted in red in Figure 1.3) represents a critical operational bottleneck.

1.2.2. The Bottleneck: Manual Image Analysis and Reporting

Zooming into the specific image analysis and reporting phase (Figure 1.3), the limitations of the current clinical pathway become evident. Currently, clinicians rely on commercial software solutions, such as *syngo.via*, to manually inspect the coronary arteries, differentiate segments, and calculate stenosis severity. Although these platforms offer semi-automated tools, the process remains highly dependent on the clinician's expertise to visually verify the vessels, identify bifurcations, and assess the plaque burden. Depending on the anatomical complexity and image quality, analyzing a single patient can consume around 30 minutes for healthy cases and 90 minutes for complex cases [1].

Once the visual assessment is complete, the extracted metrics are manually entered into an internal custom-built platform called *Filemaker*, which generates a template-based preliminary report. The spe-

cialist must then review, correct, and finalize this text before it is integrated into the *Radiology Information System* (RIS).




Figure 1.4: Detailed representation of the image analysis and reporting workflow at Hospital de Sant Pau. Adapted from Acebes Pinilla [1].

This heavy reliance on manual quantification and data entry presents significant operational challenges. Consequently, highly specialized experts spend considerable time performing repetitive, manual tasks on low-risk patients, rather than focusing their expertise on critical cases. This lack of an automated, prioritized workflow delays the diagnosis of severe cases and highlights a clear necessity for robust and automated solutions.

1.3. State of the Art: Automation of CAD Diagnosis and Visualization

The integration of Artificial Intelligence (AI) and advanced computer vision in cardiac imaging has seen exponential growth over the past decade, driven by the need to optimize clinical workflows [24]. Deep learning algorithms, particularly Convolutional Neural Networks (CNNs), have demonstrated remarkable success in medical image analysis, encouraging a move from manual CCTA assessment to automated systems [29]. However, significant limitations remain regarding their practical implementation. For these new techniques to succeed, they must be practical for hospitals, integrate smoothly into current systems, and be easily adopted by the clinical staff [22].

The prerequisite first step in any automated CAD assessment pipeline is the accurate segmentation and anatomical labeling of the coronary artery tree; before any downstream diagnostic metric can be calculated, the vessels must be correctly isolated and identified. Recent advancements have leveraged deep learning architectures, such as 3D U-Nets, to extract the vessel lumen from CCTA images [10, 34], while graph-based models like Graph Convolutional Networks (GCNs) are increasingly used to ensure topological consistency when labeling segments according to standard clinical guidelines [33]. Within Hospital de la Santa Creu i Sant Pau, preliminary research has explored these initial segmentation and labeling tasks using transformer-based models and the nn-UNet framework [11, 27].

Following successful segmentation, extracting the precise geometric properties of the vessels is essential. The Vascular Modeling Toolkit (VMTK) stands out as one of the most prominent open-source libraries for medical geometry data extraction [3]. Specifically, algorithms implemented within VMTK are widely utilized to extract accurate vascular centerlines and compute cross-sectional areas. These metrics pro-

vide the essential topological framework required to perform continuous Multi-Planar Reformations and subsequently evaluate vessel narrowing.

Once the 3D geometry of the coronary tree is established, the focus shifts to quantifying the disease burden. Current stenosis quantification techniques range from traditional intensity-based thresholding (which often struggles with severe calcifications) to modern machine learning models that predict stenosis severity directly from image features [19, 34]. Alternatively, geometric, non-ML approaches offer highly interpretable and robust methods by calculating the percentage of diameter or area stenosis based purely on the physical vessel profile, comparing the minimal lumen area against an interpolated healthy reference [18]. Recognizing the clinical value of algorithmic transparency, ongoing research at Hospital de la Santa Creu i Sant Pau has actively explored these purely geometric techniques to quantify stenosis, establishing a reliable, explainable foundation for disease assessment [9].

Beyond quantification, the effective visualization of these complex 3D metrics remains a critical challenge in the state of the art [7]. In the current clinical routine at Hospital de Sant Pau, specialists must manually navigate through cross-sectional slices and use standard viewing software to visually synthesize the data, assess the lesions, and manually fill out structured reports [1]. This cognitive load highlights a significant gap: the lack of dedicated, interactive visualization tools that bridge the gap between automated metrics and clinical reality. Modern research and clinical consensus emphasize that integrating 3D anatomical models alongside continuous luminal profiles and real medical images (such as Curved Multi-Planar Reformations) into a unified graphical interface is essential. Such integration allows clinicians to continuously validate automated findings against the ground-truth CCTA, drastically reducing the time and effort required to verify stenosis locations and seamlessly generate diagnostic reports [7].

Several commercial software solutions have emerged to address these overarching needs, including platforms like *Cleerly* [12] and *Artrya Salix* [4], which use AI to analyze arterial plaque and assess ischemia risk from CCTA images. While these robust dashboards offer highly desirable visualization components, they often function as proprietary “black boxes”. This lack of algorithmic transparency significantly hinders explainability, a feature strongly demanded by medical experts to verify how a specific CAD-RADS score, plaque distribution, or stenosis percentage was computed. Finally, these standalone commercial platforms frequently struggle to integrate well into the highly specific IT infrastructures and patient prioritization workflows of public hospitals [22].

1.4. Project Scope and Contributions

To address the clinical challenges outlined in Section 1.2, and to overcome the limitations of commercial tools discussed in Section 1.3, this project contributes directly to the comprehensive, transparent, and in-house AI-enabled diagnostic framework developed by the Dimension Lab at Hospital de la Santa Creu i Sant Pau [1]. While previous academic efforts have tackled earlier stages of this pipeline, such as patient prioritization [15] and coronary artery segmentation and labeling [11], a critical gap remains in translating raw segmented geometries into actionable clinical insights.

Therefore, the scope of this bachelor’s thesis focuses on automating the geometric quantification of coronary stenosis, predicting the associated CAD-RADS score, and integrating these findings into an interactive clinical dashboard.

Figure 1.5 provides a high-level overview of where this project sits within the pipeline, while the detailed technical workflow is reserved for the Methodology chapter.




Figure 1.5: Contextual overview of the proposed automated diagnostic framework at Hospital de la Santa Creu i Sant Pau. The contributions of this thesis (green dashed box) bridge the gap between existing hospital data systems (blue) and the final clinical report. Adapted from Acebes Pinilla [1].

By seamlessly connecting raw image analysis with the hospital’s reporting systems, this project aims to mitigate operational bottlenecks and reduce overall diagnostic times.

However, it is crucial to emphasize that the solution developed in this project is intended solely as a clinical decision-support tool, not as a replacement for medical expertise. In the highly sensitive healthcare sector, human oversight is crucial. Rather than replacing detailed diagnostic software, this visualization tool is designed to act as a rapid initial assessment interface. By presenting an immediate, high-level overview of the patient’s critical metrics, it assists specialists in quickly evaluating prioritized patient lists, helping them decide which cases require immediate, in-depth analysis. Thus, the automated findings are designed to streamline the workflow, ensuring that the final diagnostic decisions and clinical validations always remain firmly in the hands of the medical professionals.

1.5. Objectives

The primary objective of this thesis is to design and develop an automated, end-to-end pipeline that successfully bridges the current gap in the CAD diagnostic workflow. Importantly, this project does not seek to maximize the precision of individual algorithmic components; rather, the proposed functional workflow is, in itself, the core product of this work.

To achieve this overarching goal, the project addresses a sequence of intermediate objectives. First, it requires deep contextualization within the complex clinical environment to evaluate prior research conducted at the hospital and design a unified, clinically applicable framework. Building upon this foundation, the project aims to implement computational geometry algorithms to accurately extract cross-sectional areas and calculate the percentage of area stenosis from 3D models. Subsequently, a standardized logic is developed to translate these geometric metrics into a reliable CAD-RADS score prediction.

To aggregate these findings and provide cardiologists with interpretable metrics, an interactive clinical dashboard is built. Functioning as a rapid-access pop-up interface, this Support Visualization Tool gives clinicians an immediate, visual summary of the patient’s coronary status, ultimately reducing medical workload, optimizing case triaging, and significantly decreasing overall assessment time. Finally, a technical validation is conducted to verify the reliability of the extracted metrics and ensure the overall stability of the complete framework.

Chapter 2

METHODOLOGY

2.1. System Overview and Core Automated Workflow

This section formalizes the modular Python pipeline introduced in Section 1.4, translating the clinical decision-support framework into an executable computational workflow. For each patient identifier, a pre-segmented coronary lumen mask is processed sequentially into quantitative stenosis metrics, a CAD-RADS 2.0-oriented report, and an interactive visualization session.

The workflow comprises four automated blocks (Sections 2.3 to 2.6) executed as a linear chain, with each stage consuming the artifacts produced by its predecessor and no manual intervention between steps. A unified patient identifier routes ASOCA, MACS-18, and synthetic cohorts through the same implementations, while runtimes and extraction metrics are logged for Section 2.7.

2.1.1. End-to-End Data Flow




Figure 2.1: End-to-end automated workflow for a single patient case. Each block constitutes a self-contained processing stage; arrows indicate mandatory sequential dependencies and the principal data passed between stages.

Figure 2.1 summarizes the four-block architecture. The workflow assumes as input a pre-segmented coronary lumen mask and, when available, anatomical segment labels from prior segmentation and labeling work at the hospital [11].

Table 2.1: Overview of the four pipeline blocks: processing role, primary inputs, and principal outputs per patient case.

Block	Primary input(s)	Principal output(s)
1 Anatomy Extraction	Binary lumen mask	Centerline geometry, branch paths, and vessel surface models
2 Stenosis Quantification	Centerline and surface representations from Block 1	Cross-sectional area profiles, percentage area stenosis (%AS), and merged branch-level data
3 Labeling & CAD-RADS	Geometric and stenosis tables	Segment stenosis summary and patient CAD-RADS report
4 Visualization	Aggregated metrics and geometry from Blocks 1–3	Interactive clinical decision-support dashboard

To ensure explainability, the system retains intermediate mesh and tabular representations for potential clinical audits. The clinical cohorts utilized, along with the methodological depth of each processing block, are detailed below.

2.2. ASOCA and MACS-18: Coronary Segmentation and Labeling Data

In medical image analysis, *segmentation* denotes the partition of a three-dimensional scan into meaningful regions. For each voxel in the volume, an expert or algorithm decides whether it belongs to a structure of interest, such as the blood-filled interior (*lumen*) of the coronary arteries, rather than to background tissue, myocardium, or surrounding anatomy. The result is not a single measurement but a spatial map that localizes the vessel tree within the patient-specific CCTA grid. The ASOCA and MACS-18 resources provide this map, together with companion anatomical label volumes, for a shared cohort of 40 CCTA cases (20 normal and 20 diseased anatomy). These data served as the reference on which Blocks 1–4 were designed and implemented. Once the pipeline was operational, the full cohort was processed to assess end-to-end behavior. The experimental protocol and result analysis are reported in Section 2.7.

2.2.1. Segmentation and Label Volumes

For each patient case, two aligned three-dimensional grids are provided, both registered to the physical coordinate system of the source CCTA so that scale and orientation are preserved.

The *segmentation annotation* is stored as a voxel grid in NRRD format. Each voxel indicates whether it belongs to the coronary lumen or to background tissue. The pipeline treats this volume as a binary mask.

The *segment label volume* is a separate grid, stored as a compressed NIfTI file, aligned with the segmentation. Instead of a binary lumen flag, each foreground voxel stores an integer class identifier that maps to a standardized coronary segment under the AHA 17-segment model [5]. Background voxels are assigned zero.

2.2.2. Datasets: ASOCA and MACS-18

The ASOCA (*Automated Segmentation of Coronary Arteries*) challenge established one of the first large, standardized benchmarks for fully automatic segmentation of complete normal and diseased coronary trees from CCTA [17]. The resource addresses the difficulty of imaging small, moving vessels. ASOCA is publicly available and provided the original lumen masks for the 40-patient cohort. In this project, those masks served as the working reference during pipeline development and remain the default segmentation supplied to Block 1.

MACS-18 (*Multiclass Anatomical Coronary Segmentation*) is a re-segmentation of the same 40 ASOCA patient volumes, currently in development at Hospital de la Santa Creu i Sant Pau and supervised by César Acebes Pinilla [1, 2]. Unlike ASOCA, this resource is not yet publicly available. Widely used public datasets have helped train automatic segmentation methods, but their lumen masks are often evaluated mainly by how well voxels match a reference, not by whether the vessel tree forms a continuous, anatomically plausible path [2]. When the lumen mask contains breaks or fails to follow the full length of the branches, extracting reliable centerlines and quantifying vessel geometry in Blocks 1 and 2 becomes more difficult. Expert manual refinements in MACS-18 target continuous lumen connectivity from the aortic root to the most distal visible branches, yielding more anatomically faithful masks without altering the underlying imaging data, patient identity, or labeling scheme. In this project, the same AHA segment label volumes are used for both ASOCA and MACS-18 cases; only the lumen mask differs between the two sources. The pipeline implementation is otherwise identical for both datasets.

2.3. Block 1: 3D Artery Anatomy Extraction

Block 1 transforms the binary lumen mask into robust 3D vessel surfaces and mathematical centerlines. By extracting global RCA and LCA geometries alongside isolated paths from the ostium to each distal tip, it establishes the anatomical framework required for stenosis quantification in Block 2.




Figure 2.2: Block 1 hybrid workflow: spatial separation of the right and left coronary trees, automated seed placement on each tree, and extraction of centerline geometry with branch packaging.

As illustrated in Figure 2.2, this block implements a hybrid approach. Instead of relying on a single tool, it combines automated voxel analysis with the surface and centerline extraction algorithms of the Vascular Modeling Toolkit (VMTK). This design completely removes the need for manual seed-point selection while maintaining the precise 3D shape of the smoothed vessels.

Phase 1: Spatial RCA/LCA Separation

The process begins by reading the binary lumen mask using `vmtkImageReader` to preserve voxel spacing and physical origin. To process the Right Coronary Artery (RCA) and Left Coronary Artery (LCA) independently, `scipy.ndimage.label` is applied to isolate the two largest connected components. In order to correctly classify these isolated components into the main RCA and LCA trees, the algorithm calculates their spatial center of mass. This metric represents the average 3D coordinate of all voxels within each object. By sorting these center points along the patient’s spatial x -axis, the system follows standard anatomical conventions: the rightmost component is classified as the RCA, and the leftmost as the LCA. This splitting step produces two independent binary sub-volumes. Processing them separately is logically necessary before geometric extraction, as VMTK algorithms require a single starting source per connected surface.

Phase 2: Per-Artery Surface Reconstruction and Automated Seeding

The two independent binary sub-volumes outputted by Phase 1 serve as the direct input for the surface reconstruction. For each isolated tree, a 3D surface mesh is generated using `vmtkMarchingCubes`. This function creates a continuous geometric boundary formed by connected triangles by interpolating the exact boundary between the foreground lumen voxels and the background. The raw mesh is then refined via `vmtkSurfaceSmoothing` using a pass-band filter. In this context, a pass-band filter works by separating high-frequency noise from the actual anatomical structure. It removes the rough, blocky edges caused by the original voxels and this creates a clean, mathematically stable surface for the next calculations without altering the true size of the artery.

The next step requires placing seed points, which are the explicit starting and ending spatial coordinates required by VMTK to trace the centerlines. To fully automate this placement and eliminate manual user interaction, the system analyzes the overall 3D shape of the voxel mask to automatically find key structural landmarks. The algorithm to obtain these seed points works by iteratively “peeling” away the outer layers of the voxel grid (a process known as skeletonization) and was leveraged from prior work [11]. The volume is thinned using `skimage.morphology.skeletonize`, reducing the complex 3D shape down to a one-voxel-wide medial axis. Topological endpoints are then easily identified as voxels that have exactly one neighboring voxel (a node degree of one).

The ostium, where the artery leaves the aorta and centerline tracing must begin, is selected among the skeleton endpoints using the AHA label volume (Section 2.2.1). Under the AHA 17-segment model, segment 1 identifies the proximal RCA and segment 5 the proximal LCA [5]; an endpoint inside either region is designated as the ostium point. If none qualifies, the widest endpoint by maximum inscribed sphere radius is used as a fallback.

All remaining endpoints are designated as distal targets. Because these endpoints are initially identified in a discrete voxel grid (i, j, k) , they must be mapped to the continuous physical space (x, y, z) to be compatible with the mesh. This is achieved using the image origin (O_x, O_y, O_z) and the voxel spacing (S_x, S_y, S_z) through the following spatial transformation:

$$x = O_x + i S_x, \quad y = O_y + j S_y, \quad z = O_z + k S_z. \quad (2.1)$$

Once mapped, these physical coordinates are projected onto the nearest vertex of the smoothed mesh to ensure the seed points lie exactly on the mathematical manifold.

Phase 3: Centerline Extraction, Topology Labeling, and Branch Packaging

With the seed points mapped onto the mathematical manifold, `vmkCenterlines` is executed on the smoothed surface mesh. By providing the identified ostium as the source and the distal targets as the endpoints, the algorithm extracts the centerlines for each component of the coronary tree. These centerlines are 1D mathematical curves tracing exactly through the middle of the 3D vessels, creating a connected network that accurately maps the complex branching structure of the arteries. At every discrete point along these curves, the algorithm also calculates the `MaximumInscribedSphereRadius`, providing an initial estimation of the local vessel width.

Once the centerlines are generated, a labeling step assigns a specific topological point type to every coordinate based on its structural position. To guarantee anatomical validity for downstream stenosis quantification, the centerlines are partitioned into continuous branch paths running directly from the ostium to each distal endpoint. Geometrical artifacts are then removed by applying a strict filtering logic: paths containing fewer than 20 points, originating more than 4 mm from the true ostium, or failing to terminate at a valid endpoint are automatically discarded.

Finally, the extracted data is packaged into a structured format. The spatial coordinates (P_x, P_y, P_z), local radius, assigned point type, and corresponding AHA segment identifiers are stored as ordered tables for both the individual branches and the global coronary tree. The generated surface and centerline meshes are also saved in `.vtp` format for visual rendering. Concluding Block 1, the RCA and LCA tables are concatenated into a unified patient record, providing the exact geometric dataset required for Block 2.

2.4. Block 2: Geometric Stenosis Quantification

Block 2 transforms the Block 1 centerline and surface representations into percentage area stenosis (%AS) profiles by measuring cross-sectional lumen area and comparing each point against a locally interpolated healthy reference.




Figure 2.3: Block 2 workflow: orthogonal cross-sectional area extraction on each artery centerline and surface mesh, tabular integration, proximal/distal reference estimation, and %AS quantification.

As illustrated in Figure 2.3, Block 2 processes each available artery (RCA and LCA) independently. For every tree, it consumes the paired centerline and smoothed surface meshes produced in Section 2.3.

The computational logic for stenosis quantification and the methodology used to identify local healthy reference values build upon prior internal research by Burrull [9], originally grounded in the standard percentage diameter stenosis (%DS) framework described by Hideo-Kajita et al. [18]. Rather than reproducing diameter-based scores directly, this pipeline adapts those reference principles to cross-sectional *area* and percentage area stenosis (%AS).

In coronary arteries, atherosclerotic plaque frequently develops eccentrically rather than as a perfect concentric ring. Relying solely on one-dimensional diameter estimates, such as the Maximum Inscribed Sphere Radius (MISR), can therefore misrepresent luminal narrowing. In an elliptical or irregularly narrowed vessel, the inscribed sphere captures only the narrowest physical constraint and ignores the lateral space through which blood may still flow.

Conversely, the cross-sectional area captures the absolute reduction of the luminal opening regardless of cross-sectional shape. This yields a more robust representation of flow restriction and a geometric metric that correlates more closely with hemodynamic pressure drop than diameter alone [18]. The following phases mirror the diagram, detailing execution from localized area computation through patient-level aggregation.

Phase 1: Cross-Sectional Area Extraction

The process begins by lightly smoothing the global centerline and resampling it to establish a uniform spatial resolution. The polyline is interpolated so that the physical distance between consecutive discrete points is standardized to 0.1 mm by default. While larger intervals could reduce computational load, this spacing is deliberately chosen to capture high-fidelity geometric detail and subtle anatomical variations along the vessel. For exceptionally long arterial trees, this spacing is adaptively increased to ensure the total number of evaluated points does not exceed a maximum limit of 12,000 per component.

At every resampled centerline point, `vmtkCenterlineSections` computes the cross-sectional lumen area at that specific point. To do this, the algorithm finds the local direction (tangent) of the centerline and creates a flat, perpendicular cutting plane. It then slices the 3D smoothed surface mesh (from Block 1) with this plane to extract the exact 2D boundary of the vessel lumen, calculating its enclosed area in mm^2 . Any incomplete or open slices are automatically discarded.

Highly tortuous or irregular geometries can occasionally cause the native VMTK sectioning step to fail, typically as an internal kernel error. To guarantee a complete, continuous area profile, a robust fallback replicates the same orthogonal slicing logic with Visualization Toolkit (VTK) operators when a failure is detected. At each evaluation point along the centerline, a `vtkPlane` (oriented perpendicularly to the centerline) combined with a `vtkCutter` slices the 3D surface to extract the cross-sectional lumen contour and compute its enclosed area in mm^2 using a standard polygon formula. This ensures accurate cross-sectional areas even in the most challenging anatomical segments.

Phase 2: Tabular Area Integration and Branch Mapping

To perform localized stenosis quantification, the cross-sectional areas computed across the global coronary tree must be accurately mapped to the individual ostium-to-endpoint branch dataframes generated in Block 1.

For every spatial coordinate (P_x, P_y, P_z) in a branch table, if the sequence of centerline points aligns perfectly with the global area output, the value is assigned via direct row matching. If the points do

not align exactly, the system uses a `scipy.spatial.cKDTree` nearest-neighbor query to find the closest evaluated station in 3D space and map its computed area to the specific branch point.

Ultimately, this mapping step propagates a continuous Area metric to every discrete coordinate across all branch paths. It seamlessly merges the newly extracted geometry with the anatomical information inherited from Block 1, preparing the exact dataset required for the next baseline reference calculations.

Phase 3: Geodesic Distance and Reference Area Estimation

Stenosis severity is evaluated independently along each individual branch path rather than on the global coronary tree dataset. This isolation is critical: it ensures that spatial evaluation strictly follows the natural, antegrade blood flow from the ostium to the distal tip. Evaluating the tree globally could inadvertently select reference points that are spatially close in Euclidean space but belong to entirely different, non-connected vessels.

To measure distance along the tortuous 3D vessel, consecutive centerline points are connected by straight segments, and the cumulative sum of these segment lengths defines a one-dimensional geodesic distance (g_d). This g_d value is then integrated into the branch dataframe for every discrete coordinate.

To evaluate stenosis at any specific point, the algorithm requires a healthy reference area for comparison. The system first identifies proximal and distal centerline points located at geodesic positions $g_d - W$ and $g_d + W$, respectively. The reference half-window W is a crucial hyperparameter in the pipeline. Based on the geometric stenosis framework introduced by Burrull [9] and standard angiographic principles described by Hideo-Kajita et al. [18], W is set to 10 mm in this project. However, this 10 mm threshold is an empirical approximation: there is no universally optimal value, as the ideal reference distance depends on patient-specific anatomy and individual lesion length. Ideally, to maximize clinical accuracy, the system would allow a physician to manually select or adjust proximal and distal reference points based on visual anatomical assessment. Because the computed stenosis varies significantly with the automated window size, this introduces an important methodological limitation that is discussed later.

Once the two reference points are identified along the same branch, their corresponding mapped areas (A_{prox} and A_{dist}) are retrieved. The healthy reference area (A_{ref}) at the evaluation point is then calculated as the arithmetic mean of these bilateral values:

$$A_{\text{ref}} = \frac{A_{\text{prox}} + A_{\text{dist}}}{2}. \quad (2.2)$$

Because this formula requires both proximal and distal context, percentage area stenosis cannot be calculated for points at the very extremities of branch paths. Consequently, any evaluation point lying within distance W from the ostium or the distal tip is excluded, as it would possess only a one-sided reference. Alongside sensitivity to window size, additional anatomical challenges of this automated baseline (such as unpredictable behavior near major bifurcations) are addressed in Chapter 4.

Phase 4: Percentage Area Stenosis and Patient-Level Aggregation

For each qualifying centerline point, the pipeline quantifies stenosis severity. The percentage area stenosis (%AS) formula, adapted from Burrull [9], compares the local cross-sectional lumen area (A) directly against its corresponding healthy reference (A_{ref}):

$$\%AS = \left(1 - \frac{A}{A_{\text{ref}}}\right) \times 100. \quad (2.3)$$

Next, the separate branch tables are merged into a unified patient record. Since all extracted branches trace continuously from the ostium to a distal endpoint, they inherently share identical coordinates along

common arterial trunks. When these tables are combined, spatial deduplication is applied based on rounded (P_x, P_y, P_z) tuples. If multiple overlapping points exist at the same coordinate, the algorithm strictly retains the row with the highest %AS value. This conservative rule ensures that no maximum stenosis is hidden, preventing clinical underestimation of lesion severity where anatomical paths overlap.

Ultimately, Block 2 yields two structured outputs: the enriched individual branch datasets (containing g_d , A_{prox} , A_{dist} , A_{ref} , and %AS, among others) and the deduplicated global coronary tree dataset. These comprehensive tables constitute the definitive quantitative input consumed by Block 3.

2.5. Block 3: Artery Labeling and CAD-RADS 2.0 Scoring

Block 3 translates the quantitative stenosis metrics produced in Block 2 into a standardized clinical classification using the Coronary Artery Disease Reporting and Data System (CAD-RADS 2.0) framework [14].




Figure 2.4: Block 3 workflow: segment aggregation based on the AHA 17-segment model and automated CAD-RADS 2.0 scoring.

As illustrated in Figure 2.4, this block ingests the merged patient-level stenosis table from Block 2 together with the anatomical segment labels inherited from Block 1. The pipeline first summarizes the stenosis severity at the individual segment level. These segments are subsequently grouped into the primary coronary territories (RCA, LM, LAD, and LCx), establishing the essential anatomical framework required for the CAD-RADS 2.0 evaluation. Finally, the system applies rule-based logic to derive the overall patient-level CAD-RADS category and calculates the Segment Involvement Score (SIS), structuring these findings into the definitive outputs consumed by Block 4.

2.5.1. Anatomical Context: AHA 17-Segment Model

Reliable CAD-RADS assignment requires mapping each evaluated centerline point to a recognized coronary territory. The pipeline adopts the 17-segment American Heart Association (AHA) model, a widely used reporting standard for regional coronary anatomy [5]. This choice is not only clinically established

but also technically aligned with the voxel labeling scheme of the primary cohorts used in this thesis: integer labels in the ASOCA annotation volumes [17] are consumed directly as segment ID values without intermediate remapping. This direct compatibility preserves pipeline robustness and keeps segment-level aggregation consistent across patients.

Beyond individual anatomical segments, this model maps every branch into one of four primary coronary territories: the Right Coronary Artery (RCA), Left Main (LM), Left Anterior Descending (LAD), and Left Circumflex (LCX). This territorial grouping is a fundamental requirement of the CAD-RADS 2.0 scoring framework, which assigns the patient-level category from the most severe stenosis observed across these vessels [14]. Table 2.2 summarizes the complete segment-to-territory dictionary used throughout Block 3 and the clinical dashboard. Background voxels are encoded as segment 0 and are subsequently excluded from stenosis aggregation.

Table 2.2: AHA 17-segment dictionary used for anatomical labeling and CAD-RADS territory grouping. Adapted from Austen et al. [5].

ID	Abbreviation	Anatomical name	Territory
1	pRCA	Proximal RCA	RCA
2	mRCA	Mid RCA	RCA
3	dRCA	Distal RCA	RCA
4	rPDA	Right PDA	RCA
5	LM	Left Main	LM
6	pLAD	Proximal LAD	LAD
7	mLAD	Mid LAD	LAD
8	dLAD	Distal LAD	LAD
9	D1	First diagonal	LAD
10	D2	Second diagonal	LAD
11	pLCX	Proximal LCX	LCX
12	OM1	First obtuse marginal	LCX
13	dLCX	Distal LCX	LCX
14	OM2	Second obtuse marginal	LCX
15	lPDA	Left coronary PDA	LCX
16	L-ILB	LCX inferolateral branch	LCX
17	L-PLB	LCX posterolateral branch	LCX

2.5.2. Segment-Level Stenosis Aggregation

For each patient, all centerline rows in the Block 2 merged table that carry a valid `Segment_ID` are grouped by segment identifier using tabular aggregation. Within each segment, the algorithm retains the absolute maximum %AS, ensuring that focal lesions are not underestimated when multiple samples exist along the same anatomical region. Each segment is then assigned a discrete stenosis severity grade aligned with the CAD-RADS luminal stenosis bins (see Table 2.3).

Any `Segment_ID` present in the input data but outside the 1–17 dictionary (for example, labels introduced by extended annotation schemes) is flagged as *unmapped*. These rows may still appear in exploratory outputs, but they are excluded from the standard territory groupings used for CAD-RADS 2.0 triage, preventing execution failure while making non-standard labels explicit.

2.5.3. Patient-Level CAD-RADS 2.0 Classification

Patient-level scoring follows the CAD-RADS 2.0 decision principles [14]. The algorithm first determines the highest segment-level stenosis grade from the aggregated table. This grade is derived from the maximum %AS observed across all evaluated segments and subsequently classified into the standard luminal categories detailed in Table 2.3. Total coronary occlusion (category 5) is assigned when any segment reaches 100% %AS.

Categories 4A and 4B distinguish high-risk patterns among severe stenoses. As established by the CAD-RADS 2.0 framework [14], category 4B is assigned when either (i) left-main (segment 5) stenosis is at least 50% %AS, or (ii) obstructive disease ($\geq 70\%$ %AS) is present in each of the three major territories (RCA, LAD, and LCX) defined in Table 2.4. If severe stenosis is present but neither of these criteria is met, the patient is classified as CAD-RADS 4A. Categories 0–3 are assigned directly from the highest segment grade when no category 4 or 5 rule overrides the result.

Table 2.3: CAD-RADS 2.0 luminal stenosis categories applied to maximum segment-level %AS. Adapted from Cury et al. [14].

Category	Max. %AS (segment)	Clinical interpretation (summary)
0	0%	No visible stenosis
1	1–24%	Minimal stenosis
2	25–49%	Mild stenosis
3	50–69%	Moderate stenosis
4A	70–99%	Severe stenosis without CAD-RADS 4B criteria
4B	70–99% (territory pattern) or LM $\geq 50\%$	High-risk severe disease pattern
5	100%	Total coronary occlusion

Table 2.4: Major coronary territories used for CAD-RADS 2.0 triage rules in this pipeline (segment IDs follow the AHA 17-segment model). Adapted from Austen et al. and Cury et al. [5, 14].

Territory	Segment IDs	Role in automated scoring
RCA	1, 2, 3, 4	Obstructive/severe flags per vessel; contributes to three-vessel 4B rule
LM	5	Left-main $\geq 50\%$ %AS triggers CAD-RADS 4B
LAD	6, 7, 8, 9, 10	Obstructive/severe flags per vessel; contributes to three-vessel 4B rule
LCX	11, 12, 13, 14, 15, 16, 17	Obstructive/severe flags per vessel; contributes to three-vessel 4B rule

Finally, the system calculates the Segment Involvement Score (SIS), defined as the absolute number of assessable coronary segments exhibiting any measurable atherosclerotic plaque [14]. Block 3 concludes by exporting a comprehensive set of structured outputs: a segment-level stenosis summary, a patient-level CAD-RADS report (detailing the assigned category, classification rationale, highest lesion location, territory flags, and SIS), and the geometrically enriched branch tables. All these generated metrics and records are specifically formatted to serve as the direct quantitative input for the interactive clinical visualization dashboard detailed in Block 4.

2.6. Block 4: Patient Visualization Dashboard

Block 4 constitutes the final user-facing layer of the automated workflow, functioning as an interactive clinical decision-support interface built with the Streamlit framework [30]. Whereas the preceding blocks operate as fully automated processing stages, this final block translates their geometric meshes and tabular records into an explorable visual environment. Specifically, the dashboard ingests the extracted centerline coordinates, cross-sectional areas, percentage area stenosis profiles, AHA segment labels, segment-level summaries, and the patient-level CAD-RADS 2.0 report.

Conceived as a rapid evaluation pop-up rather than a standalone diagnostic system, the interface synthesizes these artifacts into a single navigable session. The design attempts to avoid “black-box” behavior, recognizing that medical professionals highly value the transparency and explainability of automated tools. Consequently, every displayed score remains directly traceable to the underlying geometric measurements.

Within the clinical workflow of Hospital de la Santa Creu i Sant Pau, this dashboard is designed to integrate seamlessly with the recently developed patient prioritization system [15]. By accessing this interface for prioritized cases, cardiologists can obtain a rapid, detailed overview to efficiently rule out low-risk patients or decide if a deeper inspection of the raw coronary computed tomography angiography (CCTA) scans is required.




Figure 2.5: Block 4 user interface layout and interaction flow: from global 3D anatomy and clinical summary to segment-level and branch-level metric validation.

As illustrated in Figure 2.5, the Streamlit dashboard is organized into four vertically stacked modules that guide the clinical review process from global anatomy to vessel-level metrics. Module 1 opens the session with interactive three-dimensional renderings of the global RCA and LCA, which can be colored by either %AS or cross-sectional area. Module 2 immediately follows with the patient clinical summary, reporting the automated CAD-RADS 2.0 category, the Segment Involvement Score (SIS), and the territory-level rationale [14]. Below these high-level overviews, Modules 3 and 4 offer dedicated segment-level and branch-level exploration. In both layers, localized three-dimensional anatomical views are shown alongside their corresponding two-dimensional metric profiles, allowing a comprehensive evaluation of the affected vessels.

Module 3: Segment Level Analysis

The third module focuses on the detailed analysis of individual segments, following the standard AHA 17-segment framework. The three-dimensional view initially displays the entire artery colored by segment identity. When a user selects a specific segment from either the RCA or LCA, the three-dimensional map highlights this selection so the user clearly sees the region under analysis. Concurrently, two synchronized two-dimensional bar plots update to display the cross-sectional area and %AS profiles. This allows the user to visualize exactly how the area and stenosis behave along the length of the chosen segment. To ensure the severity calculations are transparent and easy to verify, the charts explicitly mark the point of maximum stenosis alongside the reference areas used in Equation (2.2). This combination of three-dimensional models and two-dimensional charts is supported by visualization research [7]. The three-dimensional anatomy provides clear spatial context, while the two-dimensional profiles eliminate visual overlap, making it significantly easier to locate lesions and accurately read the metrics.

Module 4: Branch-Level Analysis

The fourth module applies the same dual-visualization principle as Module 3, but expands the scope to the continuous ostium-to-endpoint paths extracted in Block 1. Upon selecting a specific branch, the three-dimensional view highlights its full anatomical trajectory and flags the peak stenosis location. Below this, longitudinal area and %AS bar plots chart the entire branch length in a strict proximal-to-distal order. Because the stenosis quantification in Block 2 is computed independently along these paths, this view exposes the exact continuous geometric context from which the reference windows and %AS values were derived.

In Modules 3 and 4, explicitly plotting the peak stenosis alongside its bilateral reference samples guarantees full algorithmic transparency. This design allows cardiologists to rapidly validate the automated findings against their clinical intuition before deeper inspection of the raw CCTA data. Full-page screenshots demonstrating these dashboard interfaces are presented in Chapter 3.

2.7. Experimental Design and Validation Framework

Validating an automated clinical decision-support pipeline requires complementary evidence at multiple levels of abstraction. As established in the project objectives (Section 1.5), the primary goal of this thesis is to construct and integrate the complete computational workflow, rather than to maximize or optimize the clinical accuracy of stenosis quantification.

Because of anatomical limitations discussed in Section 4.2, the pipeline overestimates stenosis severity in complex regions. Consequently, the current iteration of the pipeline does not possess reliable clinical predictive value. To align with the project scope, clinical stenosis accuracy and CAD-RADS 2.0 concordance against real-world patient diagnoses are explicitly excluded from the evaluation design.

Instead, this thesis adopts a three-tier validation framework: **(i)** operational validation on real patient cohorts (ASOCA and MACS-18) to assess algorithmic robustness and geometric extraction precision; **(ii)** controlled algorithmic verification on synthetic phantoms with analytical ground truth to validate the underlying mathematical logic; and **(iii)** a pilot clinical usability study using the System Usability Scale (SUS), conducted with specialized cardiologists at Hospital de la Santa Creu i Sant Pau, to evaluate the interactive Block 4 interface. Numerical outcomes are reported in Chapter 3; interpretation of limitations and future directions is reserved for Chapter 4.

2.7.1. Real-Data Evaluation Metrics and Acceptance Criteria

The first validation tier applies a shared 40-patient cohort (20 normal and 20 diseased anatomy) defined in Section 2.2. The real-data metrics in Table 2.5 were defined to evaluate pipeline reliability, geometric extraction quality, and runtime scalability on ASOCA and MACS-18. They measure whether the workflow runs successfully, whether Block 1 recovers plausible coronary topology, and whether automated processing remains practical at cohort scale.

Runtime is reported strictly as a computational performance baseline. It is not directly comparable to the full manual CCTA review times described in Section 1.1, because this system functions as a decision-support framework rather than a standalone diagnostic tool. Clinicians are still expected to perform visual assessments; therefore, the dashboard aims to reduce the overall review burden and facilitate faster, better-informed triage when determining which patients require deeper inspection.

Table 2.5: Real-data evaluation metrics used in Sections 2.7.2 and 2.7.3.

Metric	Definition	Measurement
ESR	Execution Success Rate: fraction of cases completing Blocks 1–4 without fatal error	Automated pipeline logs
MPE	Mean Endpoints per Patient: average skeleton endpoint count per case	Block 1 output (descriptive topology richness)
EISR	Endpoint Identification Success Rate: correct endpoints / all detected endpoints	Manual visual audit per endpoint (ASOCA)
OISR	Ostium Identification Success Rate: correctly placed RCA and LCA ostium points out of 2 expected per patient	Manual visual inspection vs. aortic origin
CESR	Centerline Extraction Success Rate: successful ostium-to-endpoint paths / attempted extractions	Automated computation from Block 1 extraction records
Runtime	Wall-clock seconds per block and total per patient	Logged timestamps

Table 2.6 assigns each real-data metric a *good*, *neutral*, or *poor* band to simplify interpretation of the Chapter 3 results: *good* meets the expected level for this thesis, *neutral* denotes an acceptable but imperfect outcome, and *poor* marks a clear limitation. The thresholds were proposed by the author in consultation with clinical supervisors at Hospital de la Santa Creu i Sant Pau and should be treated as practical project benchmarks rather than official medical guidelines.

Table 2.6: Acceptance bands for real-data pipeline metrics.

Metric	Good	Neutral	Poor
ESR	$\geq 95\%$	85–94%	$< 85\%$
OISR / CESR	$\geq 85\%$	75–84%	$< 75\%$
EISR (ASOCA only)	$\geq 85\%$	75–84%	$< 75\%$
Mean Total Runtime	< 900 s	900–1800 s	> 1800 s

2.7.2. Workflow Validation on Real Datasets: ASOCA

The ASOCA resource [17] served as the development reference for Blocks 1–4. After pipeline stabilization, the full 40-patient ASOCA cohort was processed end-to-end with identical configuration to that described in Sections 2.3 to 2.6.

Goal: Evaluate geometric precision, runtime efficiency, and execution success on the initial public benchmark masks.

Procedure: For each patient, the pipeline was executed without manual intervention. The metrics defined in Table 2.5 (ESR, EISR, OISR, CESR, and per-block runtimes) were recorded for normal, diseased, and combined subgroups. Both ostium and endpoint correctness were assessed via manual visual inspection: ostium accuracy was judged by overlaying detected ostium coordinates on the 3D coronary surface and comparing them to the anatomical origins at the aortic root, while each detected endpoint was individually verified for topological validity. Finally, CESR was computed automatically from Block 1 extraction records. Comparative numerical results for this experiment are reported in Section 3.1.

2.7.3. Workflow Validation on Real Datasets: MACS-18

The identical pipeline was re-executed on the expert-refined MACS-18 lumen masks for the same 40 patient identifiers [1]. Only the input segmentation differed; segment label volumes and pipeline hyperparameters were held constant.

Goal: Determine whether the improved lumen connectivity in MACS-18 affects geometric extraction robustness, topological endpoint density, and computational cost relative to ASOCA.

Procedure: To enable a direct comparison, the same 40 patients were each processed through the pipeline without manual intervention. The applicable metrics defined in Table 2.5 (ESR, MPE, OISR, CESR, and per-block runtimes) were recorded to directly contrast both cohorts. The evaluation protocol mirrored that of Section 2.7.2: ostium correctness was assessed via manual visual inspection against the anatomical origins at the aortic root, and CESR was computed automatically from Block 1 extraction records. Comparative numerical results for this paired experiment are reported in Section 3.1.

2.7.4. Geometry and Quantification Verification: Synthetic Data

Real-patient evaluation cannot provide a voxel-independent ground truth for the percentage of area stenosis (%AS). To strictly isolate the Block 2 quantification logic, two idealized vessel phantoms were programmatically generated [9] and processed through the full pipeline using the same 10 mm reference half-window as the clinical cohorts.

Goal: Validate the underlying mathematical logic and geometric accuracy of the quantification algorithms by comparing the pipeline outputs against an absolute analytical ground truth.

Phantom Design: Both masks occupy a $100 \times 100 \times 100$ voxel grid with 1.0 mm isotropic spacing. Each phantom is a capped cylinder along the Z -axis (from $Z = 15$ to $Z = 85$) with hemispherical end caps to ensure closed surfaces compatible with VMTK cross-sectioning.

- **Synthetic 1 (Healthy).** Constant radius $R = 10$ mm. Theoretical cross-sectional area $A = \pi R^2 \approx 314.16 \text{ mm}^2$ and maximum %AS = 0%.
- **Synthetic 2 (Diseased).** Same base radius with a symmetric cosine narrowing between $Z = 40$ and $Z = 60$, reaching a minimum radius $R_{\min} = 5$ mm at $Z = 50$. Theoretical maximum area is $\approx 314.16 \text{ mm}^2$, minimum area is $\approx 78.54 \text{ mm}^2$, and maximum %AS = 75%.

The stenosis radius profile follows the equation:

$$R(z) = R_{\min} + (R_{\text{base}} - R_{\min}) \frac{1 - \cos\left(\pi \frac{z - Z_{\text{mid}}}{H}\right)}{2}, \quad z \in [40, 60], \quad (2.4)$$

with base radius $R_{\text{base}} = 10$ mm, minimum radius $R_{\min} = 5$ mm, midpoint $Z_{\text{mid}} = 50$, and half-width $H = 10$ mm.

Procedure: For each phantom, the pipeline’s peak %AS and cross-sectional area samples were directly compared against the analytical ground truth. Errors are reported as absolute differences in percentage points (for %AS) and mm² (for area). Additionally, longitudinal area and %AS profiles were extracted to provide visual verification. Table 2.7 defines the acceptance thresholds for these synthetic metrics.

Table 2.7: Acceptance bands for synthetic phantom validation.

Metric	Good	Neutral	Poor	Reference basis
Max %AS absolute error	<5%	5–10%	>10%	Inter-observer variability 5–10% [16, 18, 21]
Mean area error (healthy)	<2% rel. (≈ 6 mm ²)	2–5%	>5%	Mesh discretization tolerance
Min-area error (stenosis)	<5 mm ²	5–15 mm ²	>15 mm ²	Local neck fidelity

Numerical outcomes are reported in Section 3.2.

2.7.5. User Experience and Interface Usability Evaluation

Block 4 was evaluated through a pilot System Usability Scale (SUS) study [8] with cardiologists at Hospital de la Santa Creu i Sant Pau. SUS measures perceived interface usability; it does not assess diagnostic accuracy, consistent with the decision-support framing of Section 2.6.

Participants and procedure. Three cardiologists completed a guided session on representative pipeline outputs, then responded to a 10-item Likert questionnaire (1 = strongly disagree, 5 = strongly agree) followed by an optional open feedback prompt. The full item wording is provided in Section A.1.

Scoring. Responses were converted using the standard SUS formula: odd items contribute ($score - 1$), even items contribute ($5 - score$); the sum is multiplied by 2.5 to yield a 0–100 score per participant [8]. Mean and sample standard deviation are reported across participants.

Table 2.8: SUS score interpretation bands (Bangor adjective scale).

SUS score	Grade	Interpretation
≥ 80	Excellent	High usability [6]
68–79	Good	Above-average; acceptable for release consideration [6]
50–67	OK	Improvement needed [6]
<50	Poor	Not acceptable [6]

Pilot limitation. With $n = 3$, the study is exploratory; results indicate directional usability acceptance but should not be generalized without a larger sample (Section 4.3). Quantitative SUS outcomes appear in Section 3.4; dashboard screenshots in Section 3.3. Qualitative cardiologist feedback is analyzed only in Section 4.3.

Chapter 3

RESULTS

This chapter reports the numerical and qualitative outcomes of the three-tier validation framework defined in Section 2.7. Metrics, acceptance bands, and scope exclusions are not repeated here; each section references the corresponding methodology subsection.

3.1. Pipeline Execution: ASOCA vs. MACS-18

The full 40-patient cohort was processed end-to-end on both ASOCA and MACS-18 masks (Sections 2.7.2 and 2.7.3). Evaluation was limited to operational reliability, geometric extraction quality, and runtime; clinical stenosis accuracy was excluded (Section 2.7).

3.1.1. Execution and Endpoint Topology

Both cohorts achieved an Execution Success Rate (ESR) of 100%: all 40 normal and 40 diseased cases completed Blocks 1–4 without fatal error.

Table 3.1: Endpoint summary statistics per cohort and subgroup. MPE denotes mean endpoints per patient; EISR is reported for ASOCA only (Table 2.5).

Cohort	Subgroup	n	Mean	Std	Min	Max
ASOCA	All	40	17.1	4.78	10	32
ASOCA	Normal	20	18.2	5.61	12	32
ASOCA	Diseased	20	16.0	3.58	10	22
MACS-18	All	40	27.7	10.0	12	64
MACS-18	Normal	20	30.8	11.6	17	64
MACS-18	Diseased	20	24.7	7.15	12	42

Table 3.1 shows that MACS-18 yields a higher mean endpoint count (27.7 vs. 17.1 per patient), consistent with finer lumen geometry in the refined masks. On ASOCA, the Endpoint Identification Success Rate (EISR) was 89.77% overall (89.84% normal, 89.69% diseased), placing it in the good band of Table 2.6.




Figure 3.1: Distribution of detected endpoints per patient for ASOCA and MACS-18 cohorts.

3.1.2. Ostium and Centerline Quality

Table 3.2: Ostium accuracy (OISR) and centerline precision (CESR) by cohort and subgroup (%).

Cohort	Subgroup	Ostium (%)	Centerline (%)
ASOCA	All	85.0	83.19
ASOCA	Normal	85.0	83.18
ASOCA	Diseased	85.0	83.21
MACS-18	All	73.75	95.62
MACS-18	Normal	75.0	96.87
MACS-18	Diseased	72.5	94.04

ASOCA achieved OISR = 85.0% (good) and CESR = 83.19% (neutral) (Table 2.6). MACS-18 achieved CESR = 95.62% (good) but OISR = 73.75% (poor), reflecting the trade-off between richer topology and ostium disambiguation difficulty hypothesized in Section 2.7.3.




Figure 3.2: Grouped comparison of ostium accuracy and centerline precision between ASOCA and MACS-18.

3.1.3. Qualitative Geometry




Figure 3.3: Representative centerline extraction comparison between ASOCA and MACS-18 masks for the same patient case.




Figure 3.4: Representative smoothed surface meshes derived from ASOCA and MACS-18 segmentations for the same patient case.

Figures 3.3 and 3.4 illustrate that MACS-18 surfaces preserve more distal branches and continuous lumen paths, while ASOCA meshes exhibit coarser topology with occasional surface gaps in distal segments. These visual differences align with the higher CESR on MACS-18 and the increased endpoint density reported in Table 3.1.

3.1.4. Runtime




Figure 3.5: Per-block and total pipeline runtimes for ASOCA and MACS-18 cohorts.

Mean total runtime was 346.9 s (5.78 min) on ASOCA and 576.2 s (9.60 min) on MACS-18. Both cohort means fall within the good band (<900 s) of Table 2.6. These values reflect automated batch processing time only and are not directly comparable to full manual CCTA review, as the pipeline is in-

tended to support rather than replace expert visual assessment (Section 2.7.1). Block 2 (geometric quantification) dominated runtime on both cohorts, with a higher mean on MACS-18 (357.7 s vs. 164.6 s on ASOCA), consistent with denser geometry and more cross-section evaluations. Mean Block 1, Block 3, and Block 4 runtimes were 145.8 s, 35.2 s, and 0.16 s (ASOCA) versus 159.8 s, 57.9 s, and 0.02 s (MACS-18), respectively.

3.2. Geometric and Algorithmic Accuracy via Synthetic Data

Two analytical phantoms (Section 2.7.4) were processed through the full pipeline to verify Block 2 quantification against known ground truth.

3.2.1. Phantom Overview




Figure 3.6: Synthetic healthy and stenosis tube phantoms with extracted centerlines (Equation (2.4)).

3.2.2. Synthetic_1: Healthy Control

Table 3.3: Synthetic_1 validation metrics ($W = 10$ mm).

Case	Metric	Theoretical	Predicted	Error	Unit
Healthy tube	Max %AS	0.0	0.01	0.01	%
Healthy tube	Mean area	314.16	316.50	2.34	mm ²




Figure 3.7: Longitudinal area and %AS profiles along the healthy synthetic tube.




Figure 3.8: Three-dimensional area and %AS colormaps for the healthy synthetic tube.

The healthy phantom yielded a maximum %AS error of 0.01 percentage points and a mean area error of 2.34 mm² (0.74% relative), both within the good bands of Table 2.7.

3.2.3. Synthetic_2: Stenosis Control

Table 3.4: Synthetic_2 validation metrics ($W = 10$ mm).

Case	Metric	Theoretical	Predicted	Error	Unit
Stenosis tube	Max %AS	75.0	73.69	1.31	%
Stenosis tube	Max area	314.16	318.66	4.50	mm ²
Stenosis tube	Min area	78.54	82.55	4.01	mm ²




Figure 3.9: Longitudinal area and %AS profiles along the stenosis synthetic tube.




Figure 3.10: Three-dimensional area and %AS colormaps for the stenosis synthetic tube.

The stenosis phantom achieved a maximum %AS error of 1.31 percentage points (good: $<5\%$) and a minimum-area error of 4.01 mm^2 at the lesion neck (good: $<5 \text{ mm}^2$). All synthetic errors fall within the

good acceptance bands of Table 2.7, supporting the internal consistency of the quantification algorithm independent of real-patient segmentation uncertainty.

3.3. Coronary Artery Dashboard Visualization Output

Figures 3.11 to 3.16 present full-page screenshots of the Block 4 Streamlit dashboard for a representative diseased case, following the four-module layout of Figure 2.5.



Figure 3.11: Dashboard Module 1: interactive global RCA and LCA three-dimensional renderings with metric colormaps.




Figure 3.12: Dashboard Module 2: patient clinical summary with automated CAD-RADS 2.0 category, SIS, and territory rationale.




Figure 3.13: Dashboard Module 3: segment-level three-dimensional map and synchronized two-dimensional area and %AS bar plots.




Figure 3.14: Dashboard Module 3 (continued): segment selection and annotated peak stenosis with reference markers.




Figure 3.15: Dashboard Module 4: branch-level three-dimensional path highlighting and longitudinal metric profiles.




Figure 3.16: Dashboard Module 4 (continued): full-branch area and %AS profiles in proximal-to-distal order.

3.4. User Experience and Interface Usability Evaluation

Three cardiologists completed the SUS questionnaire after a guided dashboard session (Section 2.7.5). Figures 3.11 to 3.13 show the interface evaluated.

Table 3.5: Raw Likert responses (1 = strongly disagree, 5 = strongly agree) per SUS item.

Item	P1	P2	P3
Q1 (frequent use)	4	5	4
Q2 (unnecessarily complex)	1	3	1
Q3 (easy to use)	3	5	4
Q4 (need technical support)	3	3	1
Q5 (well integrated)	3	4	4
Q6 (inconsistency)	3	2	2
Q7 (quick to learn)	2	4	4
Q8 (cumbersome)	3	2	1
Q9 (diagnostic confidence)	3	3	4
Q10 (learn a lot beforehand)	3	2	1

Table 3.6: SUS score summary per participant and cohort statistics.

Participant	SUS score (0–100)
Participant 1	55.0
Participant 2	72.5
Participant 3	85.0
Mean	70.83
Std. dev. (sample)	15.07

The mean SUS score was 70.83 (SD = 15.07), placing the dashboard in the *Good* band (≥ 68) of Table 2.8 [6]. All participants rated frequent clinical use highly (Q1: 4 or 5). Two of three participants rated the system easy to use and intuitive (Q3 ≥ 4), and two rated its functions well integrated (Q5 ≥ 4). Diagnostic confidence (Q9) was neutral for one participant and positive for two, consistent with a decision-support rather than autonomous diagnostic role. Qualitative feedback from the evaluation session is discussed in Section 4.3.

Chapter 4

DISCUSSION

4.1. Analysis of the Results

4.2. Technical and Anatomical Limitations

4.3. Future Work

Chapter 5

CONCLUSIONS

Bibliography

- [1] C. Acebes Pinilla. An artificial intelligence framework for the prioritization and reporting of coronary artery disease patients in a cardiac imaging unit. Presentation, Hospital de la Santa Creu i Sant Pau, Dimension Lab, 2024.
- [2] C. Acebes Pinilla et al. MACS-18: Multiclass anatomical coronary segmentation with SCCT-18 regional labeling. Work in progress, Hospital de la Santa Creu i Sant Pau, 2026.
- [3] L. Antiga, M. Piccinelli, L. Botti, B. Ene-Iordache, A. Remuzzi, and D. A. Steinman. An image-based modeling framework for patient-specific computational hemodynamics. *Medical & Biological Engineering & Computing*, 46(11):1097–1112, 2008.
- [4] Artrya Ltd. Artrya salix coronary anatomy (SCA) – AI-driven detection of vulnerable plaque. <https://www.artrya.com/>, 2024.
- [5] W. G. Austen, J. E. Edwards, R. L. Frye, G. G. Gensini, V. L. Gott, L. S. Griffith, D. C. McGoon, M. L. Murphy, and B. B. Roe. A reporting system on patients evaluated for coronary artery disease. Report of the Ad Hoc Committee for grading of coronary artery disease, Council on Cardiovascular Surgery, American Heart Association. *Circulation*, 51(4 Suppl):5–40, 1975.
- [6] Aaron Bangor, Philip Kortum, and James Miller. Determining what individual SUS scores mean: Adding an adjective rating scale. *Journal of Usability Studies*, 4(3):114–123, 2009.
- [7] M. Borkin, K. Gajos, A. Peters, D. Mitsouras, S. Melchionna, F. Rybicki, et al. Evaluation of artery visualizations for heart disease diagnosis. *IEEE Transactions on Visualization and Computer Graphics*, 17(12):2479–2488, 2011.
- [8] John Brooke. SUS: A quick and dirty usability scale. In P. W. Jordan, B. Thomas, B. A. Weerdmeester, and I. L. McClelland, editors, *Usability Evaluation in Industry*, pages 189–194. Taylor & Francis, London, 1996.
- [9] E. Burrull. Stenosis quantification pipeline. Internal research, Hospital de la Santa Creu i Sant Pau, 2023.
- [10] Ö. Çiçek, A. Abdulkadir, S. S. Lienkamp, T. Brox, and O. Ronneberger. 3D U-Net: learning dense volumetric segmentation from sparse annotation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pages 424–432, 2016.
- [11] M. Clapers Colet. Automatic labeling of coronary artery segments: Multi-strategy development and evaluation. Bachelor’s thesis, Universitat Pompeu Fabra, 2025.
- [12] Cleerly, Inc. Cleerly – AI-driven heart disease care. <https://cleerlyhealth.com/>, 2024.
- [13] Cleveland Clinic. Coronary arteries: Anatomy and function. <https://my.clevelandclinic.org/health/body/22973-coronary-arteries>, 2024. Accessed: 2026-05-30.

- [14] R. C. Cury, S. Abbara, S. Achenbach, A. Agatston, D. S. Berman, M. J. Budoff, et al. CAD-RADS 2.0–2022 coronary artery disease-reporting and data system. *Journal of Cardiovascular Computed Tomography*, 16(6):536–557, 2022.
- [15] E. Ferrer Beltran. Enhanced prioritization and reporting for coronary artery disease diagnosis. Bachelor’s thesis, Universitat Pompeu Fabra, 2024.
- [16] H. M. Garcia-Garcia, M. A. Costa, and P. W. Serruys. Assessment of coronary artery disease by intravascular imaging. *European Heart Journal*, 35(35):2321–2330, 2014.
- [17] R. Gharleghi, D. Adikari, K. Ellenberger, S. Y. Ooi, C. Ellis, C. M. Chen, et al. Automated segmentation of normal and diseased coronary arteries—the ASOCA challenge. *Computerized Medical Imaging and Graphics*, 97:102049, 2022.
- [18] A. Hideo-Kajita, S. Wopperer, S. S. Beyene, Y. F. Meirovich, G. D. Melaku, K. O. Kuku, et al. Impact of two formulas to calculate percentage diameter stenosis of coronary lesions: from stenosis models (phantom lesion model) to actual clinical lesions. *The International Journal of Cardiovascular Imaging*, 35(12):2139–2146, 2019.
- [19] Y. Hong, F. Commandeur, S. Cadet, M. Goeller, M. K. Doris, X. Chen, et al. Deep learning-based stenosis quantification from coronary CT angiography. In *Proceedings of SPIE*, volume 10949, page 109492I, 2019.
- [20] Kaiser Permanente. Cardiology – artery with atherosclerosis – TPMG. Medical illustration, 2024.
- [21] K. Kalisz, J. Buethe, S. S. Saboo, S. Abbara, S. Halliburton, and P. Rajiah. Artifacts at cardiac CT: Physics and solutions. *RadioGraphics*, 36(7):2064–2083, 2016.
- [22] C. J. Kelly, A. Karthikesalingam, M. Suleyman, G. Corrado, and D. King. Key challenges for delivering clinical impact with artificial intelligence. *BMC Medicine*, 17(1):1–9, 2019.
- [23] J. Knuuti, W. Wijns, A. Saraste, D. Capodanno, E. Barbato, C. Funck-Brentano, et al. 2019 ESC guidelines for the diagnosis and management of chronic coronary syndromes. *European Heart Journal*, 41(3):407–477, 2020.
- [24] G. Litjens, T. Kooi, B. E. Bejnordi, A. A. A. Setio, F. Ciompi, M. Ghafoorian, et al. A survey on deep learning in medical image analysis. *Medical Image Analysis*, 42:60–88, 2017.
- [25] P. J. Lynch, Fred the Oyster, and M. Häggström. Coronary circulation in anterior view. Medical illustration, Wikimedia Commons, 2010.
- [26] P. A. McCullough. Coronary artery disease. *Clinical Journal of the American Society of Nephrology*, 2(3):611–616, 2007.
- [27] C. Sanchez Gomez. Coronary artery segmentation using transformer neural networks. Master’s thesis, Universitat de Barcelona, 2022.
- [28] C. Seiler. The human coronary collateral circulation. *European Journal of Clinical Investigation*, 40(5):465–476, 2010.
- [29] P. J. Slomka, D. Dey, A. Sitek, M. Motwani, D. S. Berman, and G. Germano. Artificial intelligence in cardiovascular imaging. *JACC: Cardiovascular Imaging*, 10(6):615–626, 2017.
- [30] Streamlit Inc. Streamlit: The fastest way to build and share data apps. <https://streamlit.io/>, 2024. Accessed: 2026-06-05.
- [31] World Health Organization. Cardiovascular diseases (CVDs). <https://www.who.int/health-topics/cardiovascular-diseases>, 2023.

- [32] B. Wu, A. Kheiwā, P. Swamy, M. A. Mamas, R. J. Tedford, M. Alasnag, et al. Clinical significance of coronary arterial dominance: a review of the literature. *Journal of the American Heart Association*, 13(9):e032851, 2024.
- [33] H. Yang et al. CPR-GCN: Conditional partial-residual graph convolutional network in automated anatomical labeling of coronary arteries. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 3802–3810, 2020.
- [34] X. Zhang, B. Zhang, and F. Zhang. Stenosis detection and quantification of coronary artery using machine learning and deep learning. *Angiology*, 75(5):405–416, 2024.

Appendix A

SUPPORTING MATERIAL

A.1. System Usability Scale Questionnaire

The following items were administered to cardiologist participants after a guided session with the Block 4 dashboard (Section 2.7.5). Each statement was rated on a five-point Likert scale (1 = strongly disagree, 5 = strongly agree). Item wording follows the standard SUS format [8].

1. I think that I would like to use this coronary viewer frequently in clinical practice.
2. I found the system unnecessarily complex.
3. I thought the Dashboard was easy to use and intuitive.
4. I think that I would need the support of a technical person to be able to use this system autonomously.
5. I found the various functions in this viewer (interactive 3D meshes, profile plots, clinical metrics) were well integrated.
6. I thought there was too much inconsistency in this system.
7. I would imagine that most medical professionals would learn to use this system very quickly.
8. I found the system very cumbersome to use.
9. I felt very confident using the system to analyze artery anatomy and CAD-RADS estimation.
10. I needed to learn a lot of things before I could get going with this system.

Optional open question.

Do you have any additional comments, suggestions for improvement, or features you would like to see in the future?

Scores were computed using the standard SUS conversion described in Section 2.7.5.